

DL LAB

CLASS:- CSAI-A

ROLL NO:- 71

PRN:- 12320206

NAME:- Vaishnavi Gosavi

ASSIGNMENT NO:-1

Title:- Implement Simple and Multiple Linear Regression to predict continuous variables.

1. Perform data preprocessing (handle missing values, feature scaling).
2. Fit a Simple Linear Regression model on a dataset (e.g., predicting house prices).
3. Extend to Multiple Linear Regression with multiple features.
4. Evaluate models using MSE, RMSE, and R^2 Score.
5. Visualize the regression line and predictions.

THEORY

Simple Linear Regression

Simple Linear Regression is used to establish a relationship between one independent variable (X) and a dependent variable (Y). The mathematical representation is:

$$Y = b_0 + b_1X + \epsilon$$

Where:

- **Y** represents the dependent variable (output or target).
- **X** is the independent variable (input feature).
- **b_0** is the intercept, indicating the value of **Y** when **X = 0**.

- b_1 is the slope, which shows how much Y changes with a one-unit increase in X .
- ϵ denotes the error term, which accounts for the difference between actual and predicted values.

For instance, in predicting house prices using only square footage, the model can be written as:

$$\text{House Price} = b_0 + b_1 (\text{Square Footage}) + \epsilon$$

Multiple Linear Regression

Multiple Linear Regression is an extension of Simple Linear Regression, where multiple independent variables influence the dependent variable. The general form is:

$$Y = b_0 + b_1X_1 + b_2X_2 + \dots + b_nX_n + \epsilon$$

Here:

- X_1, X_2, \dots, X_n represent different independent variables.
- b_1, b_2, \dots, b_n are the coefficients that determine the contribution of each independent variable to Y .

For example, in house price prediction, we can incorporate multiple factors:

- X_1 = Square Footage
- X_2 = Number of Bedrooms
- X_3 = Location Score

Thus, the model equation becomes:

$$\text{House Price} = b_0 + b_1 (\text{Square Footage}) + b_2 (\text{Bedrooms}) + b_3 (\text{Location Score}) + \epsilon$$

ALGORITHM

- Load the dataset (house prices with features like square footage, bedrooms, etc.).
- Preprocess the data (handle missing values, normalize if needed).
- Split into training and test sets.
- Train a linear regression model.
- Evaluate performance using metrics like Mean Squared Error (MSE), Root Mean Squared Error (RMSE) and R2.

CODE

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score

# Load dataset
df = pd.read_csv("Housing.csv")
print(df.head())
print("\nMissing values:\n", df.isnull().sum())

# Encode categorical variables
categorical_cols = df.select_dtypes(include=['object']).columns.tolist()
label_encoders = {}
for col in categorical_cols:
    le = LabelEncoder()
    df[col] = le.fit_transform(df[col])
    label_encoders[col] = le

# Define features and target variable
```

```

target = "price"
features = [col for col in df.columns if col != target]
X = df[features]
y = df[target]

# Split dataset
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Standardize features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Train simple linear regression model
single_feature = "area"
single_feature_idx = features.index(single_feature)
X_train_simple = X_train_scaled[:, single_feature_idx].reshape(-1, 1)
X_test_simple = X_test_scaled[:, single_feature_idx].reshape(-1, 1)

simple_model = LinearRegression()
simple_model.fit(X_train_simple, y_train)
y_pred_simple = simple_model.predict(X_test_simple)

# Train multiple linear regression model
multi_model = LinearRegression()
multi_model.fit(X_train_scaled, y_train)
y_pred_multi = multi_model.predict(X_test_scaled)

# Function to evaluate models
def evaluate_model(y_true, y_pred, model_name):
    mse = mean_squared_error(y_true, y_pred)
    rmse = np.sqrt(mse)
    r2 = r2_score(y_true, y_pred)
    print(f"\n{model_name} Model Performance:")
    print(f"Mean Squared Error (MSE): {mse:.2f}")
    print(f"Root Mean Squared Error (RMSE): {rmse:.2f}")
    print(f"R2 Score: {r2:.4f}")

evaluate_model(y_test, y_pred_simple, "Simple Linear Regression")
evaluate_model(y_test, y_pred_multi, "Multiple Linear Regression")

# Visualization
plt.figure(figsize=(10, 5))
sns.heatmap(df.corr(), annot=True, cmap="magma", fmt=".2f", linewidths=0.5)
plt.title("Feature Correlation Heatmap")
plt.show()

```

```
fig = px.scatter_3d(df, x='area', y='bedrooms', z='price', color='price',
                    title="3D Scatter Plot: Area vs Bedrooms vs Price",
                    opacity=0.8)
fig.show()
```

```
plt.figure(figsize=(10, 5))
sns.violinplot(data=df, palette="pastel")
plt.title("Violin Plot of Features (Outlier Detection)")
plt.xticks(rotation=45)
plt.show()
```

```
sns.pairplot(df, diag_kind="kde", plot_kws={'alpha':0.5, 's':10})
plt.show()
```

```
plt.figure(figsize=(10, 5))
sns.scatterplot(x=X_test_simple.flatten(), y=y_test, color='blue', label='Actual')
sns.scatterplot(x=X_test_simple.flatten(), y=y_pred_simple, color='red', label='Predicted')
plt.plot(X_test_simple.flatten(), simple_model.predict(X_test_simple), color='black', linewidth=2)
plt.xlabel(single_feature)
plt.ylabel("Price")
plt.title("Simple Linear Regression: Actual vs Predicted")
plt.legend()
plt.show()
```

```
plt.figure(figsize=(10, 5))
sns.lmplot(x=y_test, y=y_pred_multi, aspect=1.5, scatter_kws={'color': "purple"}, line_kws={'color':
"green"})
plt.xlabel("Actual Prices")
plt.ylabel("Predicted Prices")
plt.title("Multiple Linear Regression: Actual vs Predicted")
plt.show()
```



