

# **Assignment 8**

**Name:** Vaishnavi Gosavi

**Roll No:** 71

**Class:** TY CSAI-A

**Batch:** 2

**Title:** Use MNIST Fashion Dataset and create a classifier to classify fashion clothing into categories. Using CNN

## **Description:-**

- **Convolutional Neural Network (CNN)**

A Convolutional Neural Network (CNN) is a class of deep neural networks most commonly used in analyzing visual imagery. CNNs are inspired by the visual cortex of the human brain, which processes visual data hierarchically.

- **Architecture Components:**

- 1. Input Layer**

- Accepts images of shape (28x28x1) for grayscale MNIST data.

- 2. Convolutional Layer**

- Applies filters (kernels) to extract low-level features like edges, textures, and patterns.
- Operation: Convolution between input and filter.
- Output is called a feature map.

- 3. Activation Layer (ReLU)**

- Applies the ReLU (Rectified Linear Unit) function:  $f(x) = \max(0, x)$
- Introduces non-linearity and prevents vanishing gradients.

- 4. Pooling Layer (Max Pooling)**

- Reduces spatial dimensions (width and height) of feature maps.
- Helps in making the model more efficient and invariant to minor translations.
- E.g., 2x2 MaxPooling reduces  $28 \times 28 \rightarrow 14 \times 14$ .

- 5. Flatten Layer**

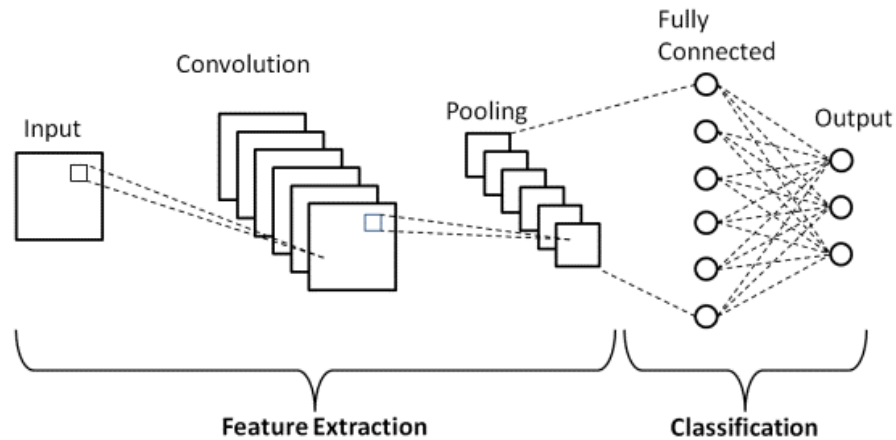
- Converts the 2D feature maps into a 1D vector to feed into the Dense (fully connected) layer.

- 6. Dense Layer**

- Fully connected layer where each neuron is connected to all outputs from the previous layer.
- Learns high-level global patterns.
- Plot actual vs. predicted stock prices.

- 7. Output Layer**

- Uses Softmax activation function to output probabilities for each of the 10 classes.
- Softmax ensures the outputs sum to 1 and represents confidence for each class.



- **Working Mechanism of CNN on MNIST**

1. Input image (28x28) is fed to the CNN.
2. First Conv layer extracts edge-level features (horizontal, vertical lines).
3. Pooling layer reduces dimensions, keeps the most important features.
4. Second Conv + Pooling layers extract more abstract features like loops, curves.
5. Flattening makes the data suitable for the Dense layer.
6. Dense layers learn complex relationships and finally classify the digit using Softmax.

- **Fashion MNIST Dataset**

- Fashion MNIST is a dataset of 28x28 grayscale images of clothing items used for multiclass image classification. It contains 70,000 images in total: 60,000 for training and 10,000 for testing, across 10 categories like T-shirt, Dress, Sneaker, and Bag.
  - Classes: 10 (e.g., T-shirt, Trouser, Coat, Sneaker)
  - Image Size: 28x28 pixels, grayscale
  - Labels: 0 to 9 (one for each class)
  - Used for: Training CNNs and benchmarking classification models



- **Algorithm**

START

1. Import necessary libraries:
  - TensorFlow, Keras, NumPy, Matplotlib (if needed)
2. Load Fashion MNIST dataset:
  - `(x_train, y_train), (x_test, y_test) = load_fashion_mnist()`
3. Preprocess data:
  - a. Normalize pixel values:
    - `x_train = x_train / 255`
    - `x_test = x_test / 255`
  - b. Reshape input data to (28, 28, 1) for CNN:
    - `x_train = reshape(x_train, (-1, 28, 28, 1))`
    - `x_test = reshape(x_test, (-1, 28, 28, 1))`
  - c. Convert labels to one-hot encoding:
    - `y_train = to_categorical(y_train, 10)`
    - `y_test = to_categorical(y_test, 10)`
4. Define CNN model:
  - a. Initialize Sequential model
  - b. Add layers:
    - `Conv2D(32 filters, 3x3 kernel, activation='relu', input_shape=(28, 28, 1))`
    - `MaxPooling2D(pool_size=(2, 2))`
    - `Conv2D(64 filters, 3x3 kernel, activation='relu')`
    - `MaxPooling2D(pool_size=(2, 2))`
    - `Flatten()`
    - `Dense(10 units, activation='softmax')`
5. Compile the model:
  - `Optimizer = 'adam'`
  - `Loss function = 'categorical_crossentropy'`
  - `Metrics = ['accuracy']`
6. Train the model:
  - `model.fit(x_train, y_train, epochs=5, batch_size=32)`
7. Evaluate the model on test data:
  - `test_loss, test_accuracy = model.evaluate(x_test, y_test)`
8. Display evaluation results:
  - Print test accuracy

END

```

import tensorflow as tf

(x_train, y_train), (x_test, y_test) =
tf.keras.datasets.fashion_mnist.load_data()
x_train = x_train.astype('float32') / 255.0
x_test = x_test.astype('float32') / 255.0
x_train = x_train.reshape(-1, 28, 28, 1)
x_test = x_test.reshape(-1, 28, 28, 1)
y_train = tf.keras.utils.to_categorical(y_train, 10)
y_test = tf.keras.utils.to_categorical(y_test, 10)

model = tf.keras.models.Sequential([
    tf.keras.layers.Conv2D(32, (3, 3), activation='relu',
input_shape=(28, 28, 1)),
    tf.keras.layers.MaxPooling2D((2, 2)),
    tf.keras.layers.Conv2D(64, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D((2, 2)),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(10, activation='softmax')
])

model.compile(optimizer='adam',
              loss='categorical_crossentropy',
              metrics=['accuracy'])

model.fit(x_train, y_train, epochs=5, batch_size=32)
loss, accuracy = model.evaluate(x_test, y_test)
print('Test accuracy:', accuracy)

```

/usr/local/lib/python3.11/dist-packages/keras/src/layers/convolutional/base\_conv.py:107: UserWarning:

Do not pass an `input\_shape`/`input\_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```

Epoch 1/5
1875/1875 ————— 55s 29ms/step - accuracy: 0.7637 -
loss: 0.6641
Epoch 2/5
1875/1875 ————— 83s 29ms/step - accuracy: 0.8774 -
loss: 0.3466
Epoch 3/5
1875/1875 ————— 81s 29ms/step - accuracy: 0.8980 -
loss: 0.2858
Epoch 4/5
1875/1875 ————— 81s 29ms/step - accuracy: 0.9044 -
loss: 0.2605
Epoch 5/5

```

```
1875/1875 _____ 81s 28ms/step - accuracy: 0.9142 -  
loss: 0.2362  
313/313 _____ 3s 8ms/step - accuracy: 0.8960 - loss:  
0.3008  
Test accuracy: 0.8950999975204468
```

```
model.summary()
```

```
Model: "sequential"
```

Layer (type) Param #	Output Shape
conv2d (Conv2D) 320	(None, 26, 26, 32)
max_pooling2d (MaxPooling2D) 0	(None, 13, 13, 32)
conv2d_1 (Conv2D) 18,496	(None, 11, 11, 64)
max_pooling2d_1 (MaxPooling2D) 0	(None, 5, 5, 64)
flatten (Flatten) 0	(None, 1600)
dense (Dense) 16,010	(None, 10)

```
Total params: 104,480 (408.13 KB)
```

```
Trainable params: 34,826 (136.04 KB)
```

```
Non-trainable params: 0 (0.00 B)
```

```
Optimizer params: 69,654 (272.09 KB)
```

```
import tensorflow as tf  
import numpy as np
```

```

from sklearn.metrics import f1_score, confusion_matrix,
accuracy_score, classification_report

y_pred = model.predict(x_test)
y_pred_classes = np.argmax(y_pred, axis=1)
y_true_classes = np.argmax(y_test, axis=1)

f1 = f1_score(y_true_classes, y_pred_classes, average='weighted')
accuracy = accuracy_score(y_true_classes, y_pred_classes)
cm = confusion_matrix(y_true_classes, y_pred_classes)
cr=classification_report(y_true_classes,y_pred_classes)
print(cr)

print("F1 Score:", f1)
print("Accuracy:", accuracy)
print("Confusion Matrix:\n", cm)

```

```

313/313 ————— 4s 13ms/step

```

	precision	recall	f1-score	support
0	0.88	0.78	0.83	1000
1	0.99	0.98	0.99	1000
2	0.88	0.82	0.85	1000
3	0.86	0.93	0.89	1000
4	0.86	0.81	0.84	1000
5	0.98	0.98	0.98	1000
6	0.65	0.79	0.71	1000
7	0.95	0.96	0.96	1000
8	0.99	0.94	0.97	1000
9	0.97	0.95	0.96	1000
accuracy				0.90
macro avg				0.90
weighted avg				0.90

F1 Score: 0.8966458262596403

Accuracy: 0.8951

Confusion Matrix:

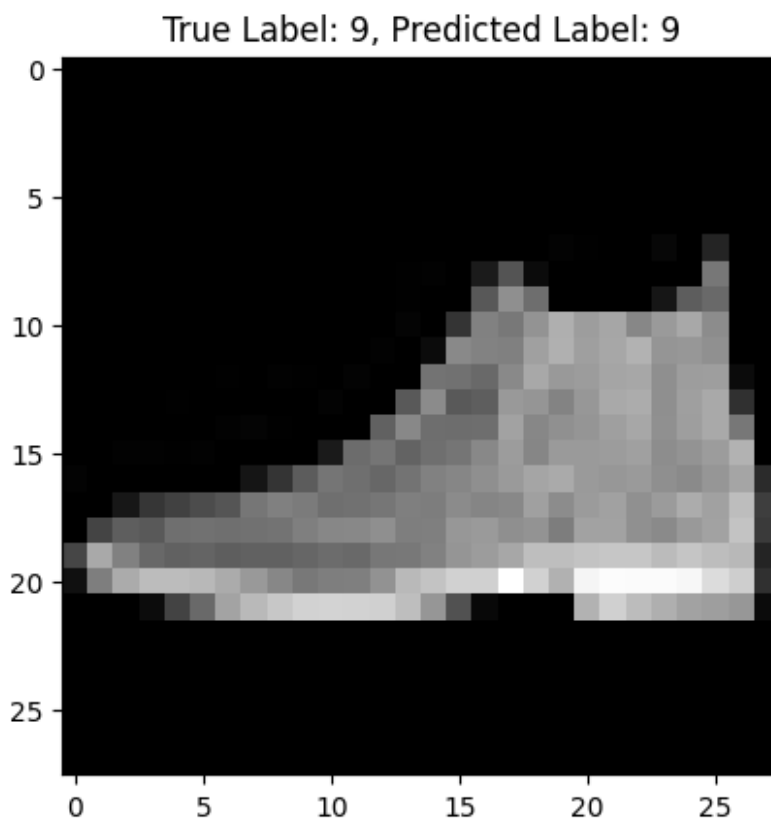
```

[[775  1  16  44  4  1 156  0  3  0]
 [ 2 979  0  16  2  0  1  0  0  0]
 [ 12  1 821  12  57  0  97  0  0  0]
 [  4  4  9 930  15  0  37  0  1  0]
 [  0  0  36  41 813  0 110  0  0  0]
 [  0  0  0  2  0 983  0  8  0  7]
 [ 77  0  48  33  50  0 788  0  4  0]
 [  0  0  0  0  0  18  0 963  1  18]
 [  5  0  8  5  1  2  32  2 944  1]
 [  1  0  0  0  0  3  0  41  0 955]]


```

```
import matplotlib.pyplot as plt
from sklearn.metrics import f1_score, confusion_matrix, accuracy_score
import numpy as np
y_pred = model.predict(x_test)
y_pred_classes = np.argmax(y_pred, axis=1)
y_true_classes = np.argmax(y_test, axis=1)
for i in range(5):
    plt.imshow(x_test[i].reshape(28, 28), cmap='gray')
    plt.title(f"True Label: {np.argmax(y_test[i])}, Predicted Label: {y_pred_classes[i]}")
    plt.show()
```


313/313 ————— 3s 9ms/step



True Label: 2, Predicted Label: 2




True Label: 1, Predicted Label: 1






True Label: 1, Predicted Label: 1



True Label: 6, Predicted Label: 6



```
feature_extractor = tf.keras.Model(inputs=model.inputs,  
outputs=model.layers[-2].output)  
features = feature_extractor.predict(x_test)
```

```
print("Feature Shape:", features.shape)
```

```
17/313 ————— 1s 7ms/step
```

```
/usr/local/lib/python3.11/dist-packages/keras/src/models/  
functional.py:237: UserWarning: The structure of `inputs` doesn't  
match the expected structure.
```

```
Expected: ['keras_tensor']
```

```
Received: inputs=Tensor(shape=(32, 28, 28, 1))
```

```
warnings.warn(msg)
```

```
313/313 ————— 2s 7ms/step
```

```
/usr/local/lib/python3.11/dist-packages/keras/src/models/  
functional.py:237: UserWarning: The structure of `inputs` doesn't  
match the expected structure.
```

```
Expected: ['keras_tensor']
```

```
Received: inputs=Tensor(shape=(None, 28, 28, 1))
```

```
warnings.warn(msg)
```

```
Feature Shape: (10000, 1600)
```

```
import numpy as np  
import pandas as pd  
from sklearn.model_selection import train_test_split  
from sklearn.metrics import classification_report  
from tensorflow.python import keras  
from tensorflow.python.keras.models import Sequential  
from tensorflow.python.keras.layers import Dense, Flatten, Conv2D,  
Dropout, MaxPooling2D  
from IPython.display import SVG  
import seaborn as sns  
import matplotlib.pyplot as plt  
%matplotlib inline  
import plotly.graph_objs as go  
import plotly.figure_factory as ff  
from plotly import tools  
from plotly.offline import download_plotlyjs, init_notebook_mode,  
plot, iplot  
init_notebook_mode(connected=True)
```

```
IMG_ROWS = 28
```

```
IMG_COLS = 28
```

```
NUM_CLASSES = 10
```

```
TEST_SIZE = 0.2
```

```
RANDOM_STATE = 2018
```

```
#Model
```

```

NO_EPOCHS = 50
BATCH_SIZE = 128
IS_LOCAL = False
import os
from tensorflow import keras
(x_train, y_train), (x_test, y_test) =
keras.datasets.fashion_mnist.load_data()
print("x_train shape:", x_train.shape)
print("y_train shape:", y_train.shape)
print("x_test shape:", x_test.shape)
print("y_test shape:", y_test.shape)

Downloading data from https://storage.googleapis.com/tensorflow/tf-
keras-datasets/train-labels-idx1-ubyte.gz
29515/29515 _____ 0s 0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-
keras-datasets/train-images-idx3-ubyte.gz
26421880/26421880 _____ 0s 0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-
keras-datasets/t10k-labels-idx1-ubyte.gz
5148/5148 _____ 0s 0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-
keras-datasets/t10k-images-idx3-ubyte.gz
4422102/4422102 _____ 0s 0us/step
x_train shape: (60000, 28, 28)
y_train shape: (60000,)
x_test shape: (10000, 28, 28)
y_test shape: (10000,)

X_train, X_test, y_train, y_test = train_test_split(
    x_train, y_train, test_size=0.3, random_state=42
)

X_train_flat = X_train.reshape(X_train.shape[0], -1) # Flatten the
images
X_test_flat = X_test.reshape(X_test.shape[0], -1)
train_data = pd.DataFrame(np.column_stack([X_train_flat, y_train]))
test_data = pd.DataFrame(np.column_stack([X_test_flat, y_test]))

X_train[0]

array([[ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0, 35, 109,
90,
        67, 95, 99, 89, 82, 61,  0,  0,  0,  0,  0,  0,
0,
        0,  0],
       [ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0, 152, 233,
244,
        226, 240, 239, 199, 135, 104,  5,  0,  0,  0,  0,  0,
0,

```

```

    178, 209, 196, 170, 180, 191, 190, 217, 164, 0, 0, 0,
0,
    0, 0],
[ 0, 0, 0, 0, 0, 0, 0, 164, 205, 186, 200, 197,
194,
    194, 205, 197, 184, 181, 188, 200, 222, 172, 0, 0, 0,
0,
    0, 0],
[ 0, 0, 0, 0, 0, 0, 0, 168, 204, 145, 158, 171,
181,
    185, 208, 196, 187, 196, 194, 186, 220, 177, 0, 0, 0,
0,
    0, 0],
[ 0, 0, 0, 0, 0, 0, 0, 154, 212, 175, 171, 171,
176,
    180, 186, 189, 187, 180, 168, 180, 215, 148, 0, 0, 0,
0,
    0, 0],
[ 0, 0, 0, 0, 0, 0, 0, 131, 243, 167, 182, 195,
194,
    189, 190, 192, 197, 197, 190, 200, 255, 186, 0, 0, 0,
0,
    0, 0],
[ 0, 0, 0, 0, 0, 0, 0, 0, 169, 156, 158, 174,
176,
    185, 190, 188, 172, 165, 149, 136, 127, 14, 0, 0, 0,
0,
    0, 0]], dtype=uint8)

```

```

import matplotlib.pyplot as plt
import numpy as np

label_indices = {}


for i in range(len(x_train)):
    label = np.argmax(y_train[i])
    if label not in label_indices:
        label_indices[label] = i
fig, axes = plt.subplots(2, 5, figsize=(10, 5))
for label, index in label_indices.items():
    ax = axes.flat[label]
    ax.imshow(x_train[index].reshape(28, 28), cmap='gray')
    ax.set_title(f"Label: {label}")
    ax.axis('off')
plt.tight_layout()
plt.show()

```

Label: 0




Label: 1



Label: 2



Label: 3




Label: 4



Label: 5




Label: 6



Label: 7



Label: 8



Label: 9

